

# Evaluating database systems relying on secure multi-party computation

Bachelor's Thesis  
in Partial Fulfillment of the Requirements for the  
Degree of  
Bachelor of Science

by  
NIKLAS ISERMANN

submitted to:  
Prof. Dr. Johannes Blömer  
and  
Prof. Dr. Juraž Somorovsky

---

Paderborn, October 15, 2022

# Eidesstattliche Versicherung

Nachname: \_\_\_\_\_ Vorname: \_\_\_\_\_

Matrikelnr.: \_\_\_\_\_ Studiengang: \_\_\_\_\_

☐ Bachelorarbeit ☐ Masterarbeit

Titel der Arbeit: Evaluating database systems relying on secure multi-party computation

☐ Die elektronische Fassung ist der Abschlussarbeit beigelegt.

☐ Die elektronische Fassung sende ich an die/den erste/n Prüfenden bzw. habe ich an die/den erste/n Prüfenden gesendet.

Ich versichere hiermit an Eides statt, dass ich die vorliegende Abschlussarbeit (Ausarbeitung inkl. Tabellen, Zeichnungen, etc.) selbstständig und ohne unzulässige fremde Hilfe erbracht habe. Ich habe keine anderen als die angegebenen Quellen und Hilfsmittel benutzt sowie wörtliche und sinngemäße Zitate kenntlich gemacht. Die Abschlussarbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen. Die elektronische Fassung entspricht der gedruckten und gebundenen Fassung.

## Belehrung

Wer vorsätzlich gegen eine die Täuschung über Prüfungsleistungen betreffende Regelung einer Hochschulprüfungsordnung verstößt, handelt ordnungswidrig. Die Ordnungswidrigkeit kann mit einer Geldbuße von bis zu 50.000,00 € geahndet werden. Zuständige Verwaltungsbehörde für die Verfolgung und Ahndung von Ordnungswidrigkeiten ist die Vizepräsidentin / der Vizepräsident für Wirtschafts- und Personalverwaltung der Universität Paderborn. Im Falle eines mehrfachen oder sonstigen schwerwiegenden Täuschungsversuches kann der Prüfling zudem exmatrikuliert werden. (§ 63 Abs. 5 Hochschulgesetz NRW in der aktuellen Fassung).

Die Universität Paderborn wird ggf. eine elektronische Überprüfung der Abschlussarbeit durchführen, um eine Täuschung festzustellen.

Ich habe die oben genannten Belehrungen gelesen und verstanden und bestätige dieses mit meiner Unterschrift.

Ort: \_\_\_\_\_ Datum: \_\_\_\_\_

Unterschrift: \_\_\_\_\_

## Datenschutzhinweis

Die o.g. Daten werden aufgrund der geltenden Prüfungsordnung (Paragraph zur Abschlussarbeit) i.V.m. § 63 Abs. 5 Hochschulgesetz NRW erhoben. Auf Grundlage der übermittelten Daten (Name, Vorname, Matrikelnummer, Studiengang, Art und Thema der Abschlussarbeit) wird bei Plagiaten bzw. Täuschung der/die Prüfende und der Prüfungsausschuss Ihres Studienganges über Konsequenzen gemäß Prüfungsordnung i.V.m. Hochschulgesetz NRW entscheiden. Die Daten werden nach Abschluss des Prüfungsverfahrens gelöscht. Eine Weiterleitung der Daten kann an die/den Prüfende/n und den Prüfungsausschuss erfolgen. Falls der Prüfungsausschuss entscheidet, eine Geldbuße zu verhängen, werden die Daten an die Vizepräsidentin für Wirtschafts- und Personalverwaltung weitergeleitet. Verantwortlich für die Verarbeitung im regulären Verfahren ist der Prüfungsausschuss Ihres Studienganges der Universität Paderborn, für die Verfolgung und Ahndung der Geldbuße ist die Vizepräsidentin für Wirtschafts- und Personalverwaltung.



# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Related Work . . . . .	2
1.2	Contribution . . . . .	3
1.3	structure . . . . .	3
<b>2</b>	<b>Preliminary</b>	<b>5</b>
2.1	Secure Multiparty Computation . . . . .	5
2.1.1	Real World and Ideal World . . . . .	5
2.2	Adversarial Models . . . . .	6
2.2.1	General Techniques . . . . .	8
2.3	Databases . . . . .	8
<b>3</b>	<b>framework description</b>	<b>9</b>
3.1	Conclave . . . . .	9
3.1.1	Optimizations . . . . .	10
3.2	ABY3 . . . . .	11
3.2.1	underlying MPC technology . . . . .	12
3.3	SMCQL . . . . .	13
3.3.1	Optimizations . . . . .	14
3.4	Discussion of Frameworks . . . . .	15
3.5	Rejected Frameworks . . . . .	17
<b>4</b>	<b>Benchmarking</b>	<b>19</b>
4.1	Measuring Runtime . . . . .	19
4.2	Networking . . . . .	20
<b>5</b>	<b>Implementation</b>	<b>25</b>
5.1	Use-Case Description . . . . .	25
5.2	Implementation . . . . .	27
<b>6</b>	<b>Evaluation</b>	<b>35</b>
6.1	use case One . . . . .	35
6.2	use case Two . . . . .	36
6.3	use case Four . . . . .	39
	<b>Bibliography</b>	<b>43</b>



# 1 Introduction

Secure multiparty computation is subfield of modern cryptography, that focuses around the question, how multiple parties with private inputs can evaluate a function over these inputs in a secure manner. In the context of multiparty computation secure means that no information about the private inputs is revealed. The first problem that was studied in the context of secure multiparty computation was Yao's "Millionaire's Problem". In Yao's "Millionaire's Problem" there are two millionaires that want to know who of them is more wealthier and do not want to disclosure any other information about their wealth. The "Millionaire's Problem" was solved by Andrew Yao in 1982 [Yao82] using Garbled Circuits.

**Secure Multiparty Computation** In the general secure multiparty computation setting(MPC), there are  $n$  parties  $P_1, \dots, P_n$  that want to compute a agreed upon public function  $F$  based on their inputs  $x_1, \dots, x_n$ . It is desired that no party  $P_i$  learns anything about the inputs of the other parties, that is not revealed by his input and the result of the function. In the example of the "Millionaire's Problem" there are two parties  $P_0, P_1$ , therefore  $n$  equals 2 and the public function is  $F = \text{argmax}(x_1, x_2)$ . In the context of the "Millionaire's Problem" secure means that  $P_1$  and  $P_1$  only learns who of  $P_1$  and  $P_2$  is more wealthier, i.e they only learn if  $x_1 > x_2$  holds or  $x_1 < x_2$  holds, as that cannot be avoided if they want to evaluate the function. There is a variety of different security models for the field of secure multiparty computation, that characterise the attacker or adversary and its capabilities. For example the attacker may be able to corrupt at most one party or possible many parties. Another important security model is the distinction between a passive adversary and an active adversary, an active adversary is allowed to force a corrupted party to change its behaviour, while a passive adversary can only observe it. For the majority of security models there already exist generic solutions that allows to evaluate an arbitrary probabilistic polynomial time function under these models. Such generic protocols like [DW82] and [Yao82] have shown to have insufficient efficiency to be applied in practice. Therefore efforts have been made to develop new more sophisticated protocols and optimise existing solutions. Archer et al. [ABL<sup>+</sup>18] and Hastings et al. [HHNZ19] have conducted an evaluation of existing multiparty computation frameworks. Yet, their studies are not exhaustive and are partly concerned with simple toy problems. Therefore, they provide only limited information about the usability of secure multiparty computation in real world applications.

**Database Applications** One use case for secure multiparty computation are large scale database systems, where multiple company's may own large databases and are willing

to corporate together and compute shared statistics over their combined data. Yet they cannot afford to disclose individual records. As these may contain highly sensitive personal information or highly valuable trade secrets. One such use case is described by Bater et al. [BEE<sup>+</sup>16], where several hospitals work together to use their patient data for scientific research. Of course, due to legal regulations, the individual patient data must remain strictly confidential. Another example for a practical utilisation of secure multiparty computation is documented by Archer et al. [ABL<sup>+</sup>18], where in a pilot project, the secure multiparty computation framework Sharemind [BJSV15] was utilized to implement a system that helped the Estonian government detect tax fraud. Since the individual data sets cannot be published, joins and aggregates operations are of particular importance. Frameworks that are included in our study are ABY3 [Rin], CipherCompute [Cos21], Conclave [VSG<sup>+</sup>19], Prio+ [AGJ<sup>+</sup>21], SMCQL [BEE<sup>+</sup>16] and VaultDB [RAB<sup>+</sup>22]. These frameworks are a very heterogeneous group and differ in their capabilities and security assumptions. For example Prio+ does support an arbitrary number of parties, while VaultDB only supports two and ABY3 works exclusively in a three party setting.

## 1.1 Related Work

**Fundamentals of Secure Multiparty Computation** As base literature in the field of MPC we have mainly relied on Cramer et al. [CD05] and Lindel [Lin17]. Cramer et al. gives a broad overview of the topic of secure multiparty computation and its theoretical foundations. Lindel provides an introduction into the topic of simulation based proof which is the key to understanding the security definitions of secure multiparty computation.

**General Purpose Compilers for Secure Multi-Party Computation** Hastings et al. [HHNZ19] did also conduct a survey about the state of secure multiparty computation. Yet there are several significant differences between their approach and ours. While Hastings et al. uses general purpose protocols that can compute arbitrary functions. We evaluate protocols that are specifically designed and optimized for the database context. Their survey was published in 2020, therefore two of the frameworks that are subject of our study, namely ABY3 and VaultDB have not been published by that time they conducted their survey. Another difference is that Hastings et al. did not collect data to measure the performance of the frameworks but rather focused on usability and documentation aspects.

**MP-SPDZ: A Versatile Framework for Multi-Party Computation** Keller et al. [HHNZ19] provide with MP-SPDZ an implementation of over 30 different MPC protocols, which enables it to compute arbitrary functions. It is notable that MP-SPDZ supports protocols for a range of different security assumptions, these include passive and active adversaries as well as settings in which the adversary may corrupt all but one party alongside with settings where the the adversary can corrupt only a single party. MP-SPDZ comes



with a high level programming interface that simplifies its usage. As MP-SPDZ is not optimized for the database operations that are of interest for us it is included in our study.

## 1.2 Contribution

In our work we have evaluated existing secure multiparty computation frameworks that are able to perform typical database operations, like joins, union, and selects. We started with a pre-selection of 6 frameworks. First, we implemented different use cases with the frameworks on a trial basis, thereby we placed particular emphasis on the usability of the frameworks. Afterwards we selected the three frameworks that were most promising to yield good results and discarded the rest. We then defined several different use cases, that include typical database operations. We implemented the use cases with the selected frameworks and evaluated their performance. Thereby we used different technical tools, to keep our results as robust and reproducible as possible. Finally, we graphically present and discuss the collected data.

## 1.3 structure

In this section we provide an overview over the rest of our work and summarise the content of each remaining chapter.

**Preliminary** In our second chapter we provide an overview over the topic of secure multiparty computation. We explain the general concept and different important security assumption. Furthermore we explain the database operations that are of importance for our work.

**Framework Description** Our third chapter contains a detailed description of the frameworks we used to perform our benchmarks. Furthermore, a comparative discussion of the frameworks can be found in the chapter, where we especially focus on usability aspects. The chapter concludes with a description of the frameworks that did not make it into our final selection.

**Benchmarking** In our fourth chapter, we describe the different metrics for which we collected data. In particular, we list the different tools that helped us to do this and provide an overview how they are used.

**Implementation** In our fourth chapter, we describe our implantation. First, we describe our use cases and motivate their design. Then we describe the concrete implementations with the individual frameworks, including commented excerpts from our source code.

**Evaluation** In our last chapter we evaluate the data we have collected. We highlight special features and provide a graphical representation of the data. Finally we conclude our work with a discussion of our results.

## 2 Preliminary

In this chapter we establish important vocabulary for our work. At first we provide an overview over the topic of Secure Multiparty Computation. That is followed by a brief explanation of some general techniques, on which the frameworks we benchmark are based on. Lastly we describe the semantics of different database operations, that are of importance for our benchmarks.

### 2.1 Secure Multiparty Computation

In the secure multiparty computation (short MPC) scenario there are  $n$  parties  $p_0, \dots, p_{n-1}$ . They want to compute an agreed upon functionality  $F(x_0, \dots, x_{n-1})$ . A functionality is a function that is allowed to have internal randomness. Each party  $p_i$  holds an input value  $x_i$ . The parties keep their input private and do not want to reveal any information about it. The goal of secure multiparty computation is to develop a protocol  $\pi$  that enables them to jointly compute  $F(x_0, \dots, x_{n-1})$ . The security goal of "not revealing the inputs" is often formalised through the Real-/Ideal-World Paradigm.

#### 2.1.1 Real World and Ideal World

When modelling security of secure multiparty computation we compare the real world, to a perfect ideal world, where the problem can be solved in a perfect way. In the perfect world, there exists a trusted third party that evaluates the function for the other parties. Because of this, no information is leaked besides the output of the function. In order for a protocol to be secure, it is required that its execution in the real world does effectively reveal no more information than the ideal world's solution.

**Real World** In the real world there exists a protocol  $\pi$  which was designed to enable the parties to compute  $F$ . All parties execute the protocol together. During the execution they communicate for several rounds. The attacker or adversary has the ability to corrupt one or more of the parties. His capability to influence the corrupted parties is an important parameter and may differ based on different security assumptions. These may range, for example, from a relatively weak adversary that can only read messages, to a very powerful adversary, which can actually influence the behaviour of the parties he corrupts. We explain the adversary models that are of importance for our benchmarks in Section 3.2 in detail. The real world view of this attacker  $A$  consists of the inputs of each corrupted party, their obtained messages throughout the protocol and the used randomness. The protocol achieves our security goal if the view of  $A$  contains no more

information beyond what can be deducted from the corrupted parties' input and outputs alone. Which we formalise by simulation in the ideal world.

**Ideal World** In the ideal world the parties can rely on a trusted, incorruptible third party  $P$  that aids them. With the aid of  $P$ , the parties can evaluate  $F$  in two simple steps. In a first round of communication every party sends  $P$  its input.  $P$  now holds all information needed to compute  $F$ . Afterwards,  $P$  sends the result to each party in a second round of communication. Like in the real world, in the ideal world there also exists an adversary. Similar to his real world counterpart, he is also able to corrupt one or more parties. Compared to his real world counterpart the ideal world adversary has otherwise very limited abilities. He can only see the input and output of the parties he corrupts. Especially the computation with the aid of  $P$  produces no intermediate results that he can observe. Depending on the underlying security assumptions he also may be able to modify the input a corrupted party sent to  $P$  in the first round of communication. For security we want to show that a real world attacker, effectively, learns nothing more than such an ideal world adversary with its clearly specified, limited capabilities. This is specified using the simulation paradigm.

**Simulation Paradigm** Given a real-world adversary  $A$ , a secure multiparty computation protocol  $\pi$ , and a functionality  $F$  for  $\pi$  to be secure we require the existence of a so called simulator  $S$ .  $S$  is an ideal world adversary for  $F$  that has to indistinguishably simulate views the real world attacker  $A$  obtains in a protocol execution. To achieve this,  $S$  is only allowed to perform corruptions consistently with  $A$ 's behaviour. This means that  $S$  and  $A$  corrupt the same parties and if  $A$  is an active adversary, then  $S$  is allowed to change the inputs of its corrupted parties. After  $S$  has performed its attack,  $S$  computes and outputs a real world view. Protocol  $\pi$  is secure against  $A$ , if the view  $S$  outputs is indistinguishable from a view of  $A$ . This means that the real world attacker  $A$  cannot learn more than the ideal world attacker  $S$ , despite the very limited abilities  $S$  has compared to  $A$ . Finally, we say  $\pi$  is secure, if, for all  $A$ ,  $\pi$  is secure against  $A$ .

## 2.2 Adversarial Models

There are multiple models and categorizations of adversaries and their capabilities. These distinctions have significant impact on feasibility and difficulty of secure multiparty computation. In the following we will outline the models and assumptions that are of importance for our benchmarks.

**Passive Adversary vs Active Adversary** A passive adversary cannot force a corrupted party to deviate from the protocol in any way. One could think of a passive adversary as a "read-only" adversary, as a passive adversary is only able to read the messages his corrupted parties receive or send. An active adversary is allowed to do everything a passive adversary is allowed, more precisely the set of all ability's an active adversary has, is a super set of the set of all abilities an passive adversary has. He has the additional

power to force a corrupted party to deviate from the protocol in an arbitrary way. So if for example, the protocol would at some point require that each party chooses an integer between 1 and  $n$  uniformly at random, then a passive adversary would have no choice but to choose the integer between 1 and  $n$  uniformly at random. On the contrary, an active adversary would be able to force a corrupted party to choose the value 42 or any other value that the adversary considers to be advantageous for him. In the ideal world a passive adversary is bound to forward the real input values. A active adversary can choose to ignore the real input and forward any value instead.

TODO motivation see 10 05

**Monolithic Adversary** A common assumption is the assumption of a so called monolithic adversary. When assuming a monolithic adversary one assumes that, there is only a single adversary that controls all corrupt parties. For the honest parties a monolithic adversary is a worst-case scenario. Because a monolithic adversary is more powerful than multiple adversaries that control the same total amount of parties but do not cooperate with each other. A protocol that is secure in the presence of a single adversary that corrupts  $n$  parties and is able to coordinate their efforts, will be secure in the presence of up to  $n$  adversary's that corrupt  $n$  parties total and do not coordinate their efforts. In the following we will assume a monolithic adversary.

**General Adversary vs Threshold Adversary** One important distinction in modelling adversaries is the distinction between a general adversary and a threshold adversary. In the threshold setting each party is a legitimate target for corruption and the threshold adversary is parametrized by a parameter  $t$ ,  $0 < t < n$ , which denotes the limit of parties it is allowed to corrupt. A common setting for  $t$  is  $t = \lfloor \frac{n}{2} \rfloor$ , which is called the honest majority. For example for  $n=3$  the presence of an honest majority means that it is assumed that the threshold adversary can corrupt at most 1 party. Threshold MPC best fits scenarios that feature a very homogenous group of parties, as it cannot model differences between parties. To model a heterogeneous group of parties one can rely on a general adversary. A general adversary is limited in his choice which party he corrupts by an adversary structure  $Z = \{Z_1, \dots, Z_l\}$ . Where  $Z_i$  can be any set of parties. The general adversary must only corrupt a set of parties  $P$  such that there exists an  $X \in Z$  that holds  $P \subset X$ . This allows for a flexible way to formalise assumptions. If, for example, in a protocol there are two parties that hold a very vital role and one wants to assume that no adversary can corrupt both of these parties, that can be formalised by using a general adversary and defining  $Z$  so that no element in  $Z$  contains both of these two parties.

**Static vs Dynamic Corruptions** Another important distinction is the distinction between static and adaptive adversaries. A static adversary is forced to choose which parties he wants to corrupt before the execution of  $\pi$  starts. An adaptive adversary can corrupt a party during the execution of the protocol. This makes the adaptive adversary much more powerful. He can try to identify "weak links" based on the information

he gets during the execution of the protocol and then choose corrupt those. Because adaptive adversaries is more powerful than a static adversary, it is harder to achieve a protocol that is secure against such an adversary. Therefore protocols that are secure against adaptive adversaries are typically significantly slower, than protocols that are only secure in the presence of a static adversary.

### **2.2.1 General Techniques**

**binary secret sharing**

**garbled circuits**

## **2.3 Databases**

coming soon

## 3 framework description

In this chapter we describe the different MPC frameworks that were candidates for our study. First we describe the frameworks for which we have implemented our use cases. We have dedicated one subchapter for each of these frameworks. First we describe the functionality they provide. Afterwards we describe the MPC technology they utilise to realise their functionality. At last we describe framework-specific details like special algorithms or other optimisations. At the end of the chapter we also provide a brief overview over some MPC frameworks that were candidates for our study but are not part of our final selection.

### 3.1 Conclave

Conclave [VSG<sup>+</sup>19] allows to perform MPC analytics on "big data". Conclave aims to provide a high-level interface that abstracts internal MPC details away from the user. Through this high abstraction level Conclave aims to make MPC more accessible for those who are not experts in this field. Every operation done with Conclave is composable, that means that the output of every query can be the input of another query. This mechanism makes it possible to construct very complex queries out of multiple relative simple queries. With Conclave one can join tables using the equivalent of an equi-join or an union operator. Conclave also supports a range of aggregate functions these include sum, mean, standard deviation.

**underlying MPC technology** Conclave utilises existing MPC frameworks for its backend to perform its underlying MPC operations. Therefore Conclave inherits most security guarantees and assumptions from these frameworks. The concrete frameworks of use are Sharemind and Obliv-C. As both of these frameworks are designed to withstand passive adversary's and do not support more than 3 parties. Conclave also assumes a passive adversary and supports up to 3 parties. Since Obliv-C is based on garbled circuits and Sharemind on secret sharing, Conclave uses both (TODO). Conclave interacts with its backend through a generic interface. Therefore it is theoretically feasible to integrate another framework to add support for more than 3 parties. Conclave assumes a passive threshold adversary that corrupts statically and is bound by an honest majority.

**Documentation** Another important property of a framework is its documentation. Conclave's documentation features three key components. First Conclave features a quick start guide. The quick start guide functions as entry point for new users and guides them through Conclave's initial setup. Second Conclave comes with an external

documentation that provides a detailed description for the majority of its functions. For each function it explains the expected behaviour as well as how input and output are supposed to be provided and received. Conclave’s external documentation archives it to provide the user an overview of its capability’s. The part of Conclave’s documentation where it shines the most are its comments. Within Conclave every important function is fully commented and even functions that are of limited significance for the end-user are partially commented. The presence of such extensive comments simplifies the usage of Conclave, as it allows to look up specify implementation details, that are missing in the external documentation, with limited effort.

### 3.1.1 Optimizations

TODO Optimisation besser beschreiben MPC techniques are multiple orders of magnitude slower then cleartext processing. Its Conclave key principle to archive better performance by avoiding the use of MPC techniques where possible. Instead of exclusively using MPC operations, Conclave evaluates queries with a combination of local cleartext processing and MPC operations. When Conclave compiles a query it applies various optimizations to it, one such optimization is conclave’s query rewriting

#### Query Rewriting    TODO

**Trust Annotations and Hybrid Operations** Conclave features optional trust annotations that allow for trade-off between security and performance. With these trust annotations one party can annotate that it does trust another party to learn the values of a specific column. There exists a variety of use cases that fit these mechanism. For example, the sensitivity of data may largely differ between columns. Therefore it may be desirable , for a party to reveal some less sensible data in order to speed up the computation. If a party decides to do so , Conclave uses these annotations to apply optimisations, that speed up query evaluation. One such optimization are Conclave’s hybrid operators. When possible Conclave substitutes expansive MPC operations with cheaper hybrid operations. In a hybrid operation one party is ”promoted” to a selectively-trusted party (short STP). Conclave reveals some input columns to the STP. Such leakage is only possible if the parties did explicitly allow it with the trust annotations. Otherwise it is not possible to apply hybrid operations. With the information the STP obtains, it can evaluate the operator using mainly local computation and only minor MPC based aid from the other parties. Besides the leakage of the input columns to the STP Conclave upholds it’s normal security guarantees for every other column. For these special operations Conclave’s security assumptions differ from its normal security assumptions and can be modelled using a general adversary. Conclave’s hybrid operations can withstand any adversary that can corrupt a set of parties that, does contain the STP but no other party, or does not contain the STP and could be withstand by a normal operation.



**Sorts and Shuffles** Many of Conclave’s high-level operators include “sub-protocols” like sorts and shuffles. These sorts and shuffles are MPC operations. As such they are highly expansive operations. Yet not all of these sorts and shuffles are always necessary. If for example a operator produces a sorted intermediate result like for example an order by operation would do, it is redundant to sort again as part of the next operator. Conclave is able to identify such redundant sorts and shuffles and eliminates them where possible. The ability to skip such expansive MPC operations provides significant performance gains.

## 3.2 ABY3

ABY3 is a 3-party MPC framework that allows to compute queries on relational database tables. It focuses on computing various SQL-like join operations as efficiently as possible. Therefore it features a large range of different join operations. These include but are not limited to left join, right join, set union, set minus, and also full joins. Besides joins it is also possible to query a single table with query’s that have a comparable semantic to the “SELECT ... FROM ... WHERE” statement in SQL. For example a selection like “select X1 from X where X2 > X1” can be done with relative ease using the implemented features of ABY3. In theory, ABY3 is able to compute any polynomial time function of a table, in practice, the efficiency may differ between functions and may not always be sufficient. For executing its MPC operations ABY3 relies mainly on secret sharing. One of ABY3 great strengths is its composability. Each operation done on one or more tables produces as output also a table, which is a valid input for another query. This allows to build larger complex applications out of many small ones, very similar like one would do with a pipes-and-filters architecture. Which is best now from the Linux command line interface(CLI).

**Prototype Implementations** ABY3 demonstrates its capability in two prototype applications. One of them illustrates the possibility of ensuring the validity of voter registration records. In the United States, each state maintains its own list of registered voters. Through the highly sensitive nature of these records coordination between states to ensure their faultlessness is not trivial. For that reason, one person moving from one state to another may often result in being registered in both states, which would allow them to illegitimacy cast a vote in both of these states. ABY3 demonstrates that it provides the states with the tools needed to detect such double registration while preserving the confidentiality of the records.

**Documentation** Probably the most common form of documentation for any software are comments within the code, in this aspect ABY3 TODO , as one on hand it features classes that are very well commented, while on the other hand there are important classes that have barely any comments in them. ABY3’s documentation also features a quick start guide that demonstrates some of its core capability’s like setting up three parties and performing basic integer operations. For external documentation ABY3 comes with

a description of how aggregate functions like MAX, SUM, COUNT can be realized when utilizing ABY3. For example, the maximum operator can be evaluated with a recursive algorithm that computes the maximum of the first and second half of the rows. Yet these descriptions focus on high level concepts and many insights that over great importance for a practical implementation remain unspecified. ABY3's pre implemented prototype implementations contribute also to its documentation, as they demonstrated how some of its more complex functions are supposed to be utilised. One aspect of documentation in which ABY3 shines is are test. ABY3 comes with 108 automated test, that significantly simplify the process of installing ABY3 and doing its initial set up, as the allow specific error tracking.

### 3.2.1 underlying MPC technology

ABY3 works within a 3 party setting with a honest majority. This is a conscious decision as the two party and tree party setting each provide their own advantages and disadvantages. The third party allows to deploy more efficient algorithms that could not be deployed in a two-party setting. For example, oblivious permutations can be done in  $O(n)$  in a three-party setting instead of  $O(n \log n)$  in a two-party setting [Rin]. ABY3 guarantees security against a passive threshold adversary. For executing its MPC operations ABY3 relies mainly on secret sharing. In order to archive composability of operations it is important that input and output have the same format. For ABY3 this requirement means that input and output of its operations are secret shared. For secret sharing based techniques having the output in a secret shared form can be archived by omitting the final reconstruction step, that would transform the output into the clear. Therefore having a secret shared based protocol be composable can be archived in a natural way. This property gives ABY3 a significant advantage over many other MPC frameworks. Pinkas Et al. [PSZ14] for example present a framework with similar capability and performance as ABY3 that is not composable as its algorithms require the input to be present in cleartext and also cannot be extended to be composable in trivial way [Rin].

ABY3's key feature are its new protocols for joins based on a MPC based cuckoo hash table. With these new protocols it is possible to join  $n$  rows with only  $O(n)$  overhead.

**Cuckoo Hashing** ABY3's version of cuckoo hashing is a special variant of a hash table. The table is based on a vector  $T$  that has  $m$  slots that can hold items. We address the  $i$ -th slot of  $T$  with  $T[i]$  for any  $i$  with  $1 \leq i \leq m$ . For a given keyspace  $X$  the hash table utilises two hash functions  $h_0, h_1: X \mapsto \{1, \dots, m\}$ . A hash function can be practically any function and is used to map the data to its storage location, as the size of the keyspace is normally larger than  $m$ , the hash functions cannot be injective. To insert any  $x \in X$  into the hash table  $x$  is inserted into  $T[h_i(x)]$  where  $i \in \{1, 2\}$  is decided with a coin toss. If  $T[h_i(x)]$  contains already an element the element is removed from the hash table and afterwards reinserted. This approach yields an important invariant(1): For any  $x$  it always holds that  $x$  is located in  $T[h_0(x)]$  or  $T[h_1(x)]$ . Therefore cuckoo hash

tables have  $O(1)$  worst case lookup time, as a lookup only needs to check two possible positions in the table.

**Computing Joins** One key task for computing any kind of join is identifying which rows have identical join keys. More precisely for two tables  $X, Y$  and key columns  $X_1, Y_1$  and any given  $i$  the question, if there exists  $j$  such that  $X_1[i] = Y_1[j]$ , must be answered, where  $X[i]$  denotes the  $i$ -th row of table  $X$  and  $X_1[i]$  the  $i$ -th entry of the first column of table  $X$ .

ABY3 implements an algorithm that solves this problem using a secure cuckoo hash table  $T$  with two hash functions  $h_1$  and  $h_2$ . In a first step each row of  $Y$  is inserted into the hash table, such that  $Y[i]$  is inserted into  $T[h_0(Y_1[i])]$  or  $T[h_1(Y_1[i])]$ . If  $X_1[i]$  has a matching join key  $Y_1[y]$ , then per definition it holds that  $X_1[i] = Y_1[y]$ , this implies also  $h_0(X_1[i]) = h_0(Y_1[y])$  and  $h_1(X_1[i]) = h_1(Y_1[y])$ . With invariant(1) the matching row  $Y[y]$ , if it exists, can only be located in  $T[h_0(X_1[i])]$  or  $T[h_1(X_1[i])]$ . Therefore in a second step a match can be found by comparing  $Y[y]$  to  $T[h_0(X_1[i])]$  and  $T[h_1(X_1[i])]$  in a secure way.

To summarize the algorithm in a more intuitive way. First the rows of  $Y$  are inserted successively into the hash table. In order to find matches the keys that are used for the hash table are the keys for the join. To check if a row  $x$  from table  $X$  has a row in table  $Y$  that matches the join criteria, one can compare it, with the two rows in the hash table that are located at the locations where one would insert  $x$ . Since hash keys and join keys are identical, every possible match of join keys is indicated by a possible collision in the hash table.

The key challenge in this algorithm is the construction and usage of a secure cuckoo hash table that does not leak sensitive information. ABY3 implements such a hash table based on an oblivious switching network [Rin].

### 3.3 SMCQL

SMCQL [BEE<sup>+</sup>16] is an MPC based framework for relational database operations that is based on an already existing MPC framework, namely OblivM [LWN<sup>+</sup>15]. With SMCQL one can specify a query and SMCQL automatically generates secure code for evaluating the query. SMCQL realizes a private data network. A private data network is a union of many mutually distrusting databases that can be queried like a single engine that holds all data of every party. From the user's perspective, a private data network functions exactly like one monolithic database. With SMCQL one can specify queries in a semantic very similar to SQL and SMCQL translates these queries into a sequence of MPC operations. Therefore SMCQL allows using MPC without having detailed knowledge of the underlying system. With this approach, SMCQL wants to increase the accessibility of MPC. SMCQL supports a variety of SQL operators, these include selection, projection, aggregation, equi-joins, theta joins, and cross products.

## Documentation

**underlying MPC technology** SMCQL currently works in a two-party setting and provides security against a passive, threshold adversary that can corrupt at most one party. So, it essentially resides in a honest majority setting. The two parties are aided by an honest broker a neutral third party that plans the execution of the protocol. Besides the honest broker planning the execution of the protocol, he is not involved in its actual execution. For this reason every MPC based operation does always include only the party's that are providing the data. The honest broker also functions as an access point for the user and receives his query. Once the honest broker receives the query it parses the query into a directed acyclic graph of operators. Each node in the graph represents one operation and an edge between two nodes annotates that the incoming node consumes the output data of the outgoing node. With the operator graph, the honest broker is able to analyze the flow of data through the query and decide which of SMCQL's different optimizations are applicable to each node. A detailed description of these optimizations can be found in Section 4.3.4. Once all optimizations are planned the honest broker generates secure MPC based code and provides it to the parties. For its secure computations SMCQL uses the already existing OblivM framework.

**Access Control** SMCQL features an access control system that enables the data owners to adequately model the sensitivity of their data. The access control is column based and each column is either public, protected, or private. A public column may always be revealed to any party including the honest broker. A protected column may be revealed if the query is  $k$ -anonymous. A query is  $k$ -anonymous if for each queried tuple it holds that the projection onto its protected attributes is indistinguishable from at least  $k-1$  other tuples. A private column is under no circumstances revealed to any party. With these access control mechanisms, SMCQL is able to speed up query evaluation by applying various optimizations. If for example an operator only works with public columns it can be evaluated without using expensive MPC operations.

### 3.3.1 Optimizations

SMCQL implements various techniques that speed up query evaluation and help it scale.

**Slicing** One such optimization is slicing. When SMCQL identifies an operator as sliceable, it partitions the input data into smaller units of computation. The partitioning of the input tuples is done horizontally. Small units of computation are easier to evaluate compared to a large monolithic operator, they allow for less complex secure code and in some cases, the evaluation of the units can be parallelized. Projections and filters are particularly easy to slice, as they can be evaluated working one tuple at a time.

**Split Operators** Another optimization that helps SMCQL scale are its split operators. A split operator splits the evaluation of an operator that requires MPC in two phases.

First, a phase of local plaintext computation that is followed by a second phase of MPC computation. The intuition behind this is that the MPC computation in the second phase is cheaper than the evaluation of the entire operator with MPC would be. Most aggregate functions can be split, in the first phase each party local aggregates over its own columns, and in the second phase, MPC is used to compute the correct aggregate out of these intermediate aggregates. The evaluation of a `count(*)` operator, for example, can be split, in the first phase each party locally counts its own input data and in the second phase these intermediate results are added up with the help of MPC.

## 3.4 Discussion of Frameworks

In this section we discuss important properties of the frameworks and compare them. At the end of the section we provide a table that summaries our discussion. In their goals and general approach Conclave and SMCQL are very similar. Both Conclave and SMCQL employ an existing MPC frameworks in a new way, with the primary goal to make MPC more accessible. Conclave even describe SMCQL as "the most similar state-of-the-art system" [VSG<sup>+</sup>19]. ABY3 on the other hand, is new framework, that devolved new MPC protocols, with the primary goal to be as much efficient as possible. Conclave, ABY3 and SMCQL all assume a passive adversary that can corrupt at most one party. Therefore their security assumptions and guarantees are quite comparable.

**Features** TODO When comparing the implemented features of the frameworks there is no It is notable that the only two features that implemented in every frameworks are the Equi-join and the UNION operators. Both Conclave and SMCQL feature a range of pre-implemented aggregate functions.

**Composability** One notable aspect of each frameworks is their approach towards composability. Both Conclave and ABY3 lay huge emphasis on it and ensure all of their operators are composable. During our work with the frameworks this mechanism has saved us a lot of time and effort, because it significantly simplified the implementation of our more complex use cases. SMCQL does not achieve full composability. On example where SMCQL fails to archive composability are select statements. As the majority of SMCQL's select statements cannot be subject to another select statement afterwards. During our work with SMCQL its limited composability has proven to impractical.

**Documentation** Comparing documentations is

**Experimental Software** One shared property of ABY3, Conclave and SMCQL, is their implementations are experimental software that primarily serves a proof of concept. This manifest especially as the existence of numerous bugs and undesirable behaviour that are part of every framework. Conclave for example features a failure detector that is supposed, to detect, if a party crashed because of an internal error and than abort the

Table 3.1: Summary ...

	ABY3	Conclave	SMCQL
2 Parties	No	Yes	Yes
3 Parties	Yes	Yes	No
Passive Adversary	Yes	Yes	Yes
Active Adversary	No	No	No
Composable	Yes	Yes	No
Quick Start Guide	Yes	Yes	No
UNION Operator	Yes	Yes	Yes
Equi-join	Yes	Yes	Yes
Left Join	Yes	No	No
Right Join	Yes	No	No
Theta Join	No	No	Yes
SUM aggregate	No	Yes	Yes
MIN/MAX aggregate	No	No	Yes
Projection	No	Yes	Yes
Automated Setup	Yes	No	No

computation. This failure detector itself is extremely error-prone and tends to produce a significant amount of false positives, i.e. it aborts the execution of the protocol, despite no party actually being crashed. Such a false positive may for example be caused, by a party that is fully occupied with a computation and does not respond to messages of the failure detector in time. For some queries this problem is more significant than for others, in extreme cases we observed more than 50% of all executions being aborted without justification. Another property that suffers from this, is the initial installation process. Installing a framework and doing its install setup is often a complex and non-trivial task. Especially for Conclave and SMCQL the instructions are often vague or incomplete. Conclave for example requires the installation of JIFF [AIL<sup>+</sup>19] a JavaScript library, as one of its dependencies. In order for JIFF to be usable for Conclave, one must actively choose to ignore several recommendations during its installation. The task is additionally complicated by the fact that, neither Conclave nor SMCQL provide sufficient test to locate errors in their setup precisely. For installing and doing its initial setup, ABY3 comes with a script that automates this task. In theory this approach simplifies the installation a lot. In practice the automated installation comes with its own issues. If such an installation does not function 100 %, fixing its errors is often a time consuming task. As its

## 3.5 Rejected Frameworks

**CipherCompute** On candidate for our study was CipherCompute [Cos21]. With the CypherCompute framework it is possible to solve a huge range of MPC problems using Rust. These include SQL operations like joins that are of interest for us. Furthermore CypherCompute provides a rich documentation, consisting of a full quickstart guide and several well documented example projects. CypherCompute utilises the SCALE-MAMBA [ACC<sup>+</sup>21] framework for its underlying MPC operations. SCALE-MAMBA itself has evolved out of the well-known SPDZ [DPSZ11] protocol. Unfortunately the early access version of CypherCompute is out of maintenance by the time we conducting this study. Therefore we have decided to not include CypherCompute in our study.

**Prio+** Prio+ [AGJ<sup>+</sup>21] is the next generation of the highly influential Prio [CGB17]. Prio+ strives to maintain the same use and security as Prio, while significantly increasing performance compared to its predecessor. Prio Plus allows an arbitrary number of parties to jointly compute aggregated statistics, like SUM, MAX, MIN operators. Prio+ utilises a client server model. In which the (potentially many) input parties use a small number of servers to compute the statistics. Prio+ guarantees confidentiality of the input values if at least one server stays honest. Unlike CipherCompute or Conclave Prio+ is not a framework for developing MPC solutions. Its rather a system for special purposes. This means that the use of Prio+ can not be extended beyond the usecases that have been originally implemented by the authors of Prio+. This leaves Prio+ with a relatively small range of usecases compared to frameworks like ABY3 or Conclave. Therefore we have decided to not include Prio+ in our study.

**VaultDB** VaultDB [RAB<sup>+</sup>22] - uses EMP toolkit [WMK16] open-source from 2016 - demonstrates proof of concept - securely computing SQL queries of private data from one or more sources - . In the US, the Health Insurance Portability and Accountability Act C++ prototype (HIPAA) of 1996 [1] defines how healthcare institutions manage, protect, and share patient data, as well as the penalties for data misuse.





## 4 Benchmarking

For benchmarking performance there is often a variety of different metrics that are relevant and require to be measured with great care, as flawed or unclean measuring may deteriorate the value of the results. In this chapter, we describe the different metrics we want to benchmark and the different tools we use to achieve clean results.

### 4.1 Measuring Runtime

A basic metric of how well a program functions, is its runtime. Time is often measured in either wall-clock time or process time. Wall-clock time references, as the name implies, the passing of time on a wall-clock while the program runs. Process time resembles the actual time a CPU was used by the program. If, for example, the process blocks for a longer period of time, its wall-clock and process execution time may differ by a large amount. Typical reasons for a for a process to block would be that it waits for a slow hard drive or for incoming network traffic from another party. One needs to be careful what time is measured and that it is measured precisely. Otherwise one may obtain flawed or unfairly biased results. We measure both wall-clock and process execution time, as this allows us a far a more detailed analysis. In particular, this approach enables us the evaluate the difference between process time and wall-clock time, which indicates how long a process was blocked. In order to do so, tool-aided measuring is required.

**Measuring Space** Besides time another metric we measure is space

**Conclave** Conclave is based on Python. Python comes with a "batteries included" approach, as it features a comprehensive standard library. Therefore plenty of tools, that provide valuable utility for us, are already integrated within python. Two such tools are `timeit` [tim21b] and the `python profiler` [cPr21], both are python libraries that offer a simple way to measure wall-clock or process execution time. `Timeit` measures exclusively end-to-end execution time, or in other words, the time the execution takes from one end to another. The `python profiler` comes with a more detailed analysis that includes detailed information on which functions have been executed, how often they have been executed, and how long it took to execute them. The extra utility provided by the `python profiler` does not come for free, compared to `timeit`, it has significant overhead that slows the execution down [cPr21]. Therefore using it would result in an unfair disadvantage for Conclave. Thus we have initially decided to use `timeit` to measure the execution time of Conclave. One significant disadvantage of `timeit`, is its inability to measure wall-clock and process time simultaneously. For this reason, in order to measure

both of these values, it is required to perform each measurement twice, which doubles the time required and is impractical. Therefore we have moved away from this approach and use time, the tool with which we measured SMCQL’s execution time.

**SMCQL** TODO time [tim21a]

**ABY3** Since ABY3 is based on C++, we can not use the python specific tools we use for Conclave. Fortunately the cryptoTools library [Rin21] is integrated into ABY3. CryptoTools is a C++ toolbelt that features a variety of tools for building cryptographic protocols. Among these utilities is a benchmarking tool for measuring runtime. With cryptoTools it is possible to measure end-to-end execution time or to measure the execution time of specific parts of the protocol. As cryptoTools provides the functionality we need and is already ”inbuilt” into ABY3, we have decided to use it for measuring ABY3’s runtime. For a detailed example of how we utilise this tool in our benchmarking, see our description of our implementation of our second use case in Section 6.

## 4.2 Networking

In our standard setup all parties run on the same machine and communicate through localhost. This simulates a practically perfect LAN connection with very low latency and high throughput. It is also of interest how well the frameworks function in less ideal conditions. Therefore we are also simulate a suboptimal wide area network(WAN) connection with high latency and limited bandwidth.

In order to do so we require a proxy server. Instead of connecting the parties to one another we connect them to the proxy server and the proxy server forwards the incoming messages to the addressed parties. To simulate a slow connection with high latency the proxy server delays incoming messages. Setting up such a proxy server is non trivial task, one challenge in particular is mapping the various connections to one another, in a correct way. This task is made more difficult by the fact that the different frameworks handle their connections in various different ways. In ABY3 for example, each party holds one direct connection to every other party. On the other hand in Conclave, every party is connected to a Node.js server that forwards the messages. In order handle these various approaches correctly, an analysis of the communication patterns is required. An additional factor that complicates our task is the fact that different frameworks use various different protocols to communicate. For example Conclave uses among others HTTP(hypertext transfer protocol),while ABY3 utilities plain TCP. Checking implementation details and source code is a target-oriented approach for such an analysis , another tool that helped us to understand the communication patterns is Wireshark.

**Wireshark** Wireshark [wir22] is an open source packet sniffer that allows to capture network traffic and saves it for a detailed analysis. Wireshark supports the analysis of numerous different protocols, among these are plain TCP, HTTP and websockets. Hence Wireshark supports all the protocols that our frameworks relies on, and therefore are

of relevance for our work. With Wireshark we have been able to record the communication of our parties and pin down the exact communication patterns. Another utility Wireshark provides for us, is the ability to record communication once our proxy was setup up. With these recordings we have been able to verify that our proxy does indeed function as intended.

**Toxiproxy** Both ABY3 and SMCQL implement communication between parties based on a plain standard socket. In the case of ABY3 it is the standard C++ socket. For SMCQL it is the standard java socket. Both of these are TCP based and can be proxied with a standard TCP proxy. For this purpose we use Shopify’s Toxiproxy [Sho22]. Toxiproxy is a Go framework that allows to simulate different hazardous network conditions. These include a connection that delays its messages to simulate a high latency setup. Once the proxy server is setup it can be configured over the command line interface (CLI) or alternatively over an HTTP(hypertext transfer protocol) interface, Toxiproxy provides multiple different dedicated HTTP clients for this purpose. The clients differ in that they offer an interface in different programming languages. We have chosen to use the provided Ruby client as it is the one that provides the most extensive documentation. A simplified example how to use Toxiproxy to simulate latency can be found in Listing 5.1. In order to use Toxiproxy, we first set up the proxy so it starts to accept new connections. This is done by calling Toxiproxy populate and specifying to which address the proxy listens to and the address the messages get forwarded to. By default Toxiproxy does not add any network limitations to a connection. In our example we apply two limitations, to simulate a latency of 1000ms. The first limitation is applied to the ”upstream” direction. Therefore it affects every message that is send towards the server from the address the server listens to. The second limitation is applied to the ”downstream” direction. Therefore it affects every response that comes from the address the server listens to.

```

0  #First we instantiate a connection between the two parties.
   Toxiproxy.populate([
2  {
   name: "aby3_party2_party1",
4  #party 3 sends its messages for party1 to port 50010 therefore the proxy
   must listen to this port
   listen: "127.0.0.1:50010",
6  #party 1 listens to port 50001 therefore theproxy must forward to this
   port
   upstream: "127.0.0.1:50001"
8  }
   ])
10 #Then we simulate a latency of 1000ms
   toxiproxy-cli add aby3_party2_party1 -type latency -name upstream latency
   =1000 -upstream
12 toxiproxy-cli add aby3_party2_party1 -type latency -name downstream
   latency=1000 -downstream
14

```

---

Listing 4.1: Setting up a proxy that simulates latency between two parties with Toxiporxy

**Node-Http-Proxy** Conclave’s communication is based partially on websockets and partially on plain standard HTTP. Websockets are implemented on top of TCP. In particular, websockets use a single TCP socket for bidirectional communication. Therefore proxying Conclave cannot be done with a simple TCP proxy. Instead we use node-http-proxy a library for proxying HTTP that also supports websockets. With node-http-proxy we are able to delay messages to simulate high latency. It is also possible to measure the amount of data sent and received over a connection. Node-http-proxy is based on JavaScript and relies heavily on an event driven programming paradigm. A simplified example how to use node-http-proxy can be found in Listing 5.2. In our example we first create a proxy server and specify the address it forwards to. In a second step we create a HTTP server that delays every incoming message for 500 milliseconds and then sends it to the proxy server. Subsequently we demonstrate how Node-Http-Proxy make use of the event driven programming paradigm. In the third step we subscribe to the “proxyReqWs” event. The “proxyReqWs” event is raised by the proxy each time an outgoing websocket message is received. The event is raised before the message is transmitted to its destination and we can submit an anonymous function as event handler. The event handler is executed for each message before the message is transmitted. We could, for example, use this mechanism to modify responses or for various other practical use cases. In our specific example, we provide a function that waits 500 milliseconds to simulate latency. In the fourth step we subscribe to the “close” event. The “close” event is raised each time a websocket is closed. We use our event handler to save the amount of bytes transmitted through the websocket. This is a showcase example, of how a proxy helps to measure important metrics that would be non trivial to measure otherwise.

```

0  #create proxy server and specify the address it forwards to
   #with the ws:true parameter we enable support for websockets
2  var proxy = new httpProxy.createProxyServer({
   target: {
4    host: 'localhost',
    port: 9005
6    },
   ws:true
8  });
   # Here we crate a standart HTTP server that delays every incoming
   message for 500ms and then forwars it the proxy server.
10 var proxyServer = http.createServer(function (req, res) {
   setTimeout(function () {
12     proxy.web(req, res);
   }, 500);
14 }).listen(9000);
   # for every outgoing message the proxy emits a proxyReqWs event that we
   react to and delays the message for 500 milliseconds
16 proxy.on('proxyReqWs', function(){

```

```
18     setTimeout(function(){ },500)
    });
    # each time a connection is closed the proxy emits a "close" event that
    # we react to and save the amount of bytes transmitted through the
    # connection
20 proxy.on('close',function (res,socket,head) {
    send = socket.bytesRead;
22     received = socket.bytesWritten;
    });
24
```

Listing 4.2: Setting up a proxy that simulates latency with node-http-proxy



## 5 Implementation

In this chapter we describe the use-cases we have chosen to implement and benchmark. We first describe and motivate our choice of use-cases and, afterwards, give reasonable details on their implementation. We do implement four use-cases that are ordered from the least complex to the most complex.

### 5.1 Use-Case Description

We have decided to implement every use-case with Conclave and SMCQL for two parties. As ABY3 does not allow a two party protocol but requires at least three parties, we have implemented everything for ABY3 with three parties. For ABY3 there will always be only two parties that provide input data. The third party will not provide input data but will assist in the execution of the protocol. We consider this approach justified, as it is the exact approach Conclave and ABY3 use in their Evaluation. And if the authors of the frameworks considers such an approach fair, then we do hold no objections against it. We will refer to first party that provides input as Alice and to the second party that provides input as Bob. Furthermore if a framework is not capable of replicating the semantic of a use-case to one hundred percent, we do use that functionality available to implement the most similar semantic the framework can provide. For example use-case two features boolean logic, but Conclave does not feature an explicit boolean logic. Therefore we use integers where the integer one represents the value true and the integer zero represents false.

**Use-Case One** For our first use-case, we have chosen a single join. A single join may not be a very complex use-case but it is not irrelevant as joins are of great importance for practically every relational database query. It is estimated that over 60 % of privacy-sensitive analytics queries include at least one join [JNS17]. In our first experiment, Alice and Bob each hold one table. Each of these tables consists of 4 columns. The first column serves as the primary key that is used for the join. The other 3 columns are filled with random non-negative integers and simulate user data. We have chosen to use non-negative integers, because the usage of negative integers has resulted in an increased frequency of various bugs and other undesirable behaviour. We are calculating an equijoin with the primary key generated in such a way that 50 % of the entries in each table will match the join criteria. We do not generate the primary key randomly, because the size of output relation depends on it and we want those sizes to be consistent between individual executions. As result, we will reveal the entire outcome of the join.

So the primary utility provided by the use of MPC operations is the fact that the entries that do not match the join criteria are obscured.

```
0 SELECT *  
FROM Alice A JOIN Bob B  
2 ON A.primary_key = B.primary_key
```

Listing 5.1: Functional equivalent SQL statement for our first use-case

**Use-Case Two** Computing joins alone is of limited use if the result of the join can not be subject to further selection. Therefore in our second use-case we will first compute a join and the query the result with a classic "SELECT ... FROM ... WHERE" statement. Similar to first experiment Alice and Bob again each hold one table. The tables consist of two columns. The first column serves as primary key and the second column contains a boolean values that is generated at random. For our experiment we will in a first step compute the natural join of the two tables, where the primary key column zips the two tables together. The primary keys will be again generated in such a way that 50 % of the columns will be included in the join. In a second step we will apply a where filter to the result of the join and eliminate every row that does not have two identical boolean values. This use case functions also as a simple showcase example for composability and will show how well this mechanism functions in practise.

```
0 SELECT PrimaryKey , AliceBool , BobBool  
FROM AliceTable  
2 NATURAL JOIN BOB TABLE  
WHERE AliceBool = BobBool
```

Listing 5.2: Functional equivalent SQL statement for our second use-case

**Use-Case Three** Besides joins, another very important group of SQL operations are aggregate functions. Over a third of all privacy-sensitive analytics queries requires a aggregation [JNS17]. Therefore our third use-case is centred around an aggregate function, or more precisely a maximum operator. TODO

**Use-Case Four** For our fourth and last use-case, we compare two special features of SMCQL and Conclave. Both Conclave and SMCQL feature a mechanic, that allows revealing some of the columns of the input data. The revelation of the input data allows them to apply optimizations that speed up computation while preserving the privacy of the other columns. For a more detailed description see SMCQL's access control and Conclave's trust annotations. Therefore in our fourth use-case, we are going to replicate the setup of our first use-case but this time we will allow the leakage of the primary key column. Such leakage will allow both Conclave and SMCQL to apply their optimization. Replicating the setup of the first use-case enables us, to compare the results of the fourth use-case to the first use-case. This comparison will show how big the speed up of these optimizations is in practice.



## 5.2 Implementation

In this section we provide a description of notable details of the implementation of our use cases. The descriptions are ordered from the least complex to the most complex use case. Since SMCQL only requires the query to be specified in a valid SQL syntax and does the rest of the implementation automatically, we focus here on Conclave and ABY3 as their implementation is more complex.

**Use Case One** As our first use case is the least complex one, we use it for a full demonstration how ABY3's and Conclave's basic workflow functions.

**Conclave** The first step in Conclave is always the definition of the relation scheme of the input. Accordingly, we first define two lists of columns, one for Alice and one for Bob. Each column is defined with a name, a data type, and an integer that indicates its owner. The owner is used to annotate trust and apply optimizations, as described in our framework description of Conclave. In our first use case each party trust itself and only itself. Once the relation scheme is defined we can populate it by using the `create` function. Create loads the data from a .CSV(Comma-Separated-Values ) file that we generated previously. Once the tables are populated we can compute the join. In order to, compute a join in conclave we need to parse the two tables of the join, a name for the result, and the two columns that are used as key columns for the join. Lastly, we can reveal the output of our computation using the `collect` function. The `collect` function requires two parameters, one that specifies which table is revealed and a second one that specifies to whom it is revealed.

```

0 def protocol():
1     # define the schema for the input tables
2     AliceColumns = [
3         defCol("primary_key", "INTEGER", [1]),
4         defCol("user_data1_Alice", "INTEGER", [1]),
5         defCol("user_data2_Alice", "INTEGER", [1]),
6         defCol("user_data3_Alice", "INTEGER", [1]),
7     ]
8     BobColumns = [
9         defCol("primary_key", "INTEGER", [2]),
10        defCol("user_data1_Bob", "INTEGER", [2]),
11        defCol("user_data2_Bob", "INTEGER", [2]),
12        defCol("user_data3_Bob", "INTEGER", [2]),
13    ]
14    # the content of the tables is loaded from a pregenerated .CSV
15    AliceTable = create("AliceTable", AliceColumns, {1})
16    BobTable = create("BobTable", BobColumns, {2})
17    # calculate the join over the two tables
18    JOIN = join(AliceTable, BobTable, 'JOIN', ['primary_key'], ['primary_key',
19    ])
20    # reveal the output of the join to Alice
21    collect(JOIN, 1)
22    # reveal the output of the join to Bob

```

```
22 collect(JOIN, 2)
```

Listing 5.3: The Python protocol of Conclave for our first use case

**ABY3** Before the actual execution of the protocol starts ABY3 requires some setup. Similar to Conclave, the first step in ABY3 is also the definition of the relation scheme. Hence we start with the definition of two lists of [?]. Each column is described by a name, a data type, and some data type-dependent information. In our case, we use integers and annotate that we use 32-bit integers. After the relation scheme is defined we create local tables and populate them. One last step that needs to be done before the execution of the protocol starts is the initialization of a message server that manages the communication between the parties. Once the execution of the protocol starts we need to convert our local input into a table that is secret shared between the parties. Before we can compute the join of the shared tables we need to define which columns of the shared tables we want to select. Therefore we define a list of references to the shared columns. In order to obtain the desired semantic, the list needs to contain each unique column and one reference to each column that is a duplicate in both tables. In our use case the columns containing the user data are unique and the key columns are duplicates because they have identical naming and data types. Therefore our list contains seven columns. In the next step we compute the join by providing the join function with one reference to each key column and a list of columns we want to select. The output of the join is still a shared table. To obtain the plaintext values of our result we need to explicitly convert the shared table.

```
0 void use_case_two(u32 rows, u32 cols) {
1     std::vector<ColumnInfo> AliceCols = { ColumnInfo{ "key", TypeID::IntID,
2         32}};
3     std::vector<ColumnInfo> BobCols = { ColumnInfo{ "key", TypeID::IntID,
4         32}};
5     for (u32 i = 1; i < cols; ++i)
6     {
7         AliceCols.emplace_back("Alice" + std::to_string(i), TypeID::IntID, 32);
8         BobCols.emplace_back("Bob" + std::to_string(i), TypeID::IntID, 32);
9     }
10    # Create tables for Alice and Bob and fill them with content
11    Table AliceTable(rows, AliceCols);
12    Table BobTable(rows, BobCols);
13    u32 intersectionsize = rows * 0.5;
14    # Fill the primary columns such that 50% of all entries match the join
15    # criteria
16    for (u64 i = 0; i < rows; ++i)
17    {
18        # if out is false then the entry will be included in the join
19        auto out = (i >= intersectionsize);
20        AliceTable.mColumns[0].mData(i, 0) = i + 1;
21        BobTable.mColumns[0].mData(i, 0) = i + 1 + (rows * out);
22    }
23    # Fill the other columns with random integers
```

```

22   for (u64 i = 1; i < cols; ++i){
23       for (u64 j = 0; j < rows; ++j){
24           AliceTable.mColumns[i].mData(j, 0) = rand();
25           BobTable.mColumns[i].mData(j, 0) = rand();
26       }
27   }
28   # Instantiate server that handels communication
29   DBServer server[3];
30   ...
31   # Here the exectution of the protocoll starts.
32   auto protocol = [&](int i) {
33       # create a secret shared version of the input
34       SharedTable AliceSharedTable = (i == 0) ?
35       server[i].localInput(AliceTable) : server[i].remoteInput(0);
36       SharedTable BobSharedTable = (i == 1) ?
37       server[i].localInput(BobTable) : server[i].remoteInput(1);
38       # define a list of clumns we want to select
39       std::vector<SharedTable::ColRef> Select_columns;
40       for (u64 i = 0; i < cols; ++i){
41           Select_columns.emplace_back(SharedTable::ColRef(BobSharedTable,
42           BobSharedTable.mColumns[i]));
43       }
44       for (u64 i = 1; i < cols; ++i){
45           Select_columns.emplace_back(SharedTable::ColRef(AliceSharedTable,
46           AliceSharedTable.mColumns[i]));
47       }
48       # compute the join
49       SharedTable JOIN = srvs[i].join( SharedTable::ColRef(AliceSharedTable,
50       AliceSharedTable.mColumns[0]), SharedTable::ColRef(BobSharedTable,
51       BobSharedTable.mColumns[0]), First_Select_columns);
52       # reveal the plaintext values of the result
53       aby3::i32Matrix result = RevealAll(server[i], JOIN);
54   };
55
56   auto AliceThread = std::thread(protocol, 0); start Alice's thread
57   auto BobThread = std::thread(protocol, 1); start Bob's thread
58   thread(protocol, 2); # start the assisting third party
59   t0.join(); # wait for Alice to finish
60   t1.join(); # wait for Bob to finish
61 }

```

Listing 5.4: Simplified Protocol for our first use case in ABY3

**use case Two** For our second use case, we do not implement the query directly. Instead, we apply an optimized query that has an identical output. To obtain the desired result Alice and Bob compute two intermediate results, the union of these intermediate results will then yield our final result, an equivalent SQL statement is listened in Listing 6.5. In the first step, Alice and Bob filter their input once so that it contains only entries with the value false in its boolean column. Afterward, they can compute the join of these two filtered tables. The result of this join is the first important intermediate result, as each of its entries is also part of the final result. The second important intermediate result

can be obtained by the same procedure when filtering the input for entries that contain true. The union of these two intermediate results is our final result.

We have chosen to implement the use case in this indirect way for two reasons. First Conclave cannot evaluate the query directly, as Conclave’s current prototype implementation can apply a WHERE filter exclusivity to a table that is either the direct input of a party or the output of a unary operator, that has only one table as input. Additionally, our indirect implementation has a significantly better performance compared to the direct implementation, for a detailed analysis see our evaluation in Chapter 7.

```

0 SELECT PrimaryKey, AliceBool, BobBool FROM
  (SELECT PrimaryKey, AliceBool FROM AliceTable WHERE AliceBool = false)
2 NATURAL JOIN
  (SELECT PrimaryKey, BobBool FROM AliceTable WHERE BobBool = false)
4 UNION
  (SELECT PrimaryKey, AliceBool, BobBool FROM
6   (SELECT PrimaryKey, AliceBool FROM AliceTable WHERE AliceBool = true)
   NATURAL JOIN
8   (SELECT PrimaryKey, BobBool FROM BobTable WHERE AliceBool = true))

```

Listing 5.5: Functional equivalent SQL statement for our optimized implementation of our second use case

**Correctness** Despite it not being obvious, our procedure indeed yields the correct result, as each join generates only entries, that have a primary key that is included in Bob’s and Alice’s inputs and have identical boolean values. Therefore every entry that is part of our result is also part of the correct result.

If any entry is not part of our result, it cannot be part of either intermediate result. Any entry that is not part of either intermediate result, features a primary key that is either not represented in both inputs or contains two boolean values that are not equal. Therefore the entry cannot be part of the correct result. Consequently, each entry that is not part of our result, cannot be part of the correct result. Which is the contraposition and therefore equivalent, to the fact that each entry of the correct result is part of our result. As our result contains every entry of the correct result and the correct result contains every entry of our result, our result, and the correct result are identical. Therefore our procedure is indeed correct.

**Conclave** In order to implement our second use case in Conclave, we start with the import of the input. Therefore we annotate the layout of each table. Conclave does not support a dedicated data type for boolean values, consequently, we use integers for every column. In the second step, we generate four new tables by applying the filters to the input. These filters are a showcase example of how Conclave can speed up computation by applying optimizations. A naive approach for applying the filters would be to share the input tables and then filter the shared tables using an MPC algorithm. This is an approach Conclave is in theory capable of. But Conclave is able to utilize the fact, that each filter has only a single input that is known to one party in the clear. In

Consequence, the parties can apply the filter to their input locally without the use of an MPC algorithm. After they have applied the filters the parties can then create shared tables of the filtered input. In the next step compute the two intermediate results by joining the shared filter input of Alice and Bob. Here Conclave demonstrates how it supports composability, as the output of the filter function is used as input for the join function. Finally, we compute the row-wise concatenation of the two intermediate results and publish the final result to Alice and Bob.

```

0 def protocol():
    AliceColumns = [
2     defCol("primary_key", "INTEGER", [1]),
    defCol("AliceBool", "INTEGER", [1])
4     ]
    AliceTable = create("AliceTable", AliceColumns, {1})
6     BobColumns = [
    defCol("primary_key", "INTEGER", [2]),
8     defCol("BobBool", "INTEGER", [2])
    ]
10    BobTable = create("input_2", AliceColumns, {2})
    AliceFilterFalse = cc_filter(AliceTable, "AliceFilterFalse", "AliceBool",
    "=", scalar=0)
12    BobFilterFalse = cc_filter(BobTable, "AliceFilterFalse", "BobBool", "=",
    scalar=0)
    AliceFilterTrue = cc_filter(AliceTable, "BobFilterTrue", "AliceBool", "="
    ", scalar=1)
14    BobFilterFalse = cc_filter(BobTable, "BobFilterTrue", "BobBool", "=",
    scalar=1)
    intermediateFalse = join(AliceFilterFalse, BobFilterFalse, '
    intermediateFalse', ['primary_key'], ['primary_key'])
16    intermediateTrue = join(AliceFilterTrue, BobFilterTrue, 'join_result1',
    ['primary_key'], ['primary_key'])
    final_result = concat([ intermediateFalse, intermediateTure ], "
    final_result" )
18    collect(final_result, 1)
    collect(final_result, 2)

```

Listing 5.6: Simplified Protocol for our second use case in Conclave

**ABY3** Our second use case is also a showcase example of how we use ABY3's integrated benchmarking capability to sample valuable information. With the in ABY3 integrated timer, we are able to set **TimePoints** at arbitrary portions of the protocol. After the execution is finished, the timer returns how much time is passed before each individual **TimePoints** was reached. For this reason, we obtain detailed information, on how the different parts of the computation compose its total runtime. Since Conclave is able to apply the filter locally, in order to hold the comparison between them fair, we also filter the locally. Therefore Alice and Bob apply the filters to their input using standard C++ before the protocol begins. The majority of the work is done by two functions we present in Listing 6.6. The first Function generates the intermediate results, in order to do so, in the first step we load the input and convert it into a secret shared table. In the next

step we calculate the join over the shared tables, to do so we need to pass a reference towards the two key columns, together with a list of the columns we want to select over. Our second function uses the first function to obtain the intermediate results. Once we have obtained those we compute their union.

```

0  # include the timer and instanciate it
   include "cryptoTools/Common/Timer.h"
2  Timer t;
   ...
4  # generate input and apply filter to input locally, omitted for reason of
   simplicity.
   ...
6  # i denotes the party id where i==0 means Alice, i==1 Bob and i==2
   denotes the third assisting party
   auto getIntermediate = [&](int i, int filter){
8      SharedTable AliceTable;
      SharedTable BobTable;
10     # Load the prefiltered input and convert it into a secret share.
       if (filter == 1) {
12         AliceTable = (i == 0) ? srvs[i].localInput(AliceInputFilterOne):
           srvs[i].remoteInput(0);
14         BobTable = (i == 1) ? srvs[i].localInput(BobInputFilterOne):
           srvs[i].remoteInput(1);
16     } else {
18         AliceTable = (i == 0) ? srvs[i].localInput(AliceInputFilterZero):
           srvs[i].remoteInput(0);
20         BobTable = (i == 1) ? srvs[i].localInput(BobInputFilterZero):
           srvs[i].remoteInput(1);
       }
       # Here we annotate that we want to select the first two columns of
       Alice's Table and the second column of Bob's Table
22     std::vector<SharedTable::ColRef> First_Select_columns;
       Select_columns.emplace_back(SharedTable::ColRef(BobTable, BobTable.
           mColumns[0]));
24     Select_columns.emplace_back(SharedTable::ColRef(BobTable, BobTable.
           mColumns[1]));
       Select_columns.emplace_back(SharedTable::ColRef(AliceTable, AliceTable.
           mColumns[1]));
26     # We calculate the join and return it
       return srvs[i].join( SharedTable::ColRef(AliceTable, AliceTable.
           mColumns[0]), SharedTable::ColRef( BobTable, BobTable.mColumns[0]),
           Select_columns);
28
   };
30 auto use_case2 = [&](int i, int filter)
   {
32     t.setTimePoint("start");
       SharedTable IntermediateOne = getIntermediate(i,0);
34     t.setTimePoint("FirstIntermediate");
       SharedTable IntermediateZero = getIntermediate(i,1);
36     t.setTimePoint("SecondIntermediate");
       SharedTable UNION;
38     # Here we annotate which columns of the first Intermediate
       # will part of the UNION

```

```

40  std::vector<SharedTable::ColRef>  SelectOnes;
41  std::vector<SharedTable::ColRef>  SelectZeros;
42  for(u64 index=0; index <3; ++index) {
43      # Here we annotate that all columns of both intermediateresults
44      # will be part of the UNION
45      SelectOnes.emplace_back(SharedTable::ColRef( IntermediateOne ,
46      IntermediateOne.mColumns[index]));
47      SelectZeros.emplace_back(SharedTable::ColRef(IntermediateZero ,
48      IntermediateZero.mColumns[index]));
49  }
50  t.setTimePoint("UNION");
51  # Compute the union of the two intermediate results
52  UNION = srvs[i].rightUnion(SharedTable::ColRef( IntermediateZero ,
53  IntermediateZero.mColumns[0]) ,
54  SharedTable::ColRef( intermediateOne , intermediateOne.mColumns[0]) ,
55  SelectZeros , SelectOnes);
56  t.setTimePoint("end");
57
58  };
59  auto t0 = std::thread(use_case2, 0,1); #Start Alice thread
60  auto t1 = std::thread(use_case2, 1,1); #Start Bob thread
61  use_case2(2,1); # Start the assisting third party
62  t0.join(); # Wait for Alice to finish
63  t1.join(); # Wait for Bob to finish
64  write_to_file(t) # Write timer information to file for further analysis

```

Listing 5.7: Simplified Protocol for our second use case in ABY3

**use case Three** coming soon sollte aber nicht so viel sein da nicht mit aby3 implementiert wird

**use case Four** As our first and fourth use case a nearly identical their implementation differs only in minor details. Especially for Conclave the implementation is identical besides the fact that we need to use different trust annotations. In order to allow Conclave to apply its public join optimization, we annotate that each party is allowed to learn both of the key columns.

```

0  def protocol():
1  # define the schema for the input tables
2  AliceColumns = [
3  # annotate that Alice trust Bob to learn hear key column
4  defCol("primary_key", "INTEGER", [1,2]), #
5  defCol("user_data1_Alice", "INTEGER", [1]),
6  defCol("user_data2_Alice", "INTEGER", [1]),
7  defCol("user_data3_Alice", "INTEGER", [1]),
8  ]
9  BobColumns = [
10 # annotate that Bob trust Alice to learn his key column
11 defCol("primary_key", "INTEGER", [1,2]),
12 defCol("user_data1_Bob", "INTEGER", [2]),
13 defCol("user_data2_Bob", "INTEGER", [2]),

```

```
14 defCol("user_data3_Bob", "INTEGER", [2])  
15 ]  
16 ...  
17 ...  
18 collect(JOIN, 1)  
collect(JOIN, 2)
```

Listing 5.8: The Python protocol of Conclave for our last use case



## 6 Evaluation

In this chapter we discuss the results of our benchmarks. We provide figures for the individual performance of the frameworks and highlight notable details. We evaluate the result of each use case individually, ordered from least complex to most complex. At the end of the chapter we discuss our results and provide our final conclusion.

**Experimental Setup** All our benchmarks are conducted on a single machine that hosts all parties. Our machine runs on a 64 bit Linux operations system, more precisely we use Debian 11 with the Linux Kernel version 5.10.0-18.amd64. The system features 4 CPUs that each run at 2,3 GHz and 128 GB of RAM. In order to fit the results of different frameworks into shared figures, we use for the axis of our figures a logarithmic scaling to base 10. We always annotate the total amount of input-rows included in a computation. For example if a computation involves two parties that each provide an input of 100 rows, we will annotate that as 200 rows of input. To obtain more reliable measurements, we perform each measurement three times, wherever possible, and select the median of our measurements as final result.

### 6.1 use case One

For our first evaluation we are inspecting the result of our first use case in the localhost setting. We have visualised a comparison of runtime and space consumption in Figure 6.1 and 6.2.

**Time** For SMCQL we have been able to evaluate the query for up to 140 input-rows. In order to compute the query with 140 input-rows SMCQL took about 10 hours. For Conclave we have been able to scale up to 500 input-rows, with 500 row taking about 8 hours. We have observed that the difference in speed between SMCQL and Conclave grows significantly larger with increasing input size. For an input of size 20, Conclave is 5 times faster than SMCQL, for an input of size 100 Conclave already is more than 50 times faster and for an input of size 140, Conclave is more than 75 times faster than SMCQL. In the first use case ABY3 is by far the best scaling framework. ABY3 is able to compute the query with 140 and 500 input columns in less than a second. ABY3 runtimes does not increase in any significant way before its input reaches the mark of 128000 input-rows, for which it requires 1107 milliseconds. We have been able to evaluate the query with 16.000.000 input-rows in one minute and 46 seconds.




Figure 6.1: Runtime of ABY3, Conclave and SMCQL for our first use case

**Space** Its notable that for large input sizes SMCQL is more space efficient then Conclave, while for small input input sizes Conclave is more space efficient then Conclave. They break even at by an input size of 60 rows, for which they both require about 870 MB of space. For input larger then 60 rows SMCQL scales significantly better then Conclave. For an input of size 140 rows SMCQL only needs about 1300MB while Conclave more then 4 times that much. As the space consumption of SMCQL grows in a very linear fashion and the space consumption of Conclave doubles in relatively short intervals, we can only assumes that this trend would continuo for large datasets. Similar to the case of the runtime, ABY3 also performs significantly better then Conclave and SMCQL. For input size 140, ABY3 only allocates 28MB of space and therefore is over 46 times more efficient then SMCQL. For input size 500, ABY3 allocates about 32 MB of space and is therefore 1800 times more efficient then Conclave, which requires over 58 GB of space.

## 6.2 use case Two

In chapter 5 we have described an optimized way to evaluate the query of our second use case. In order to see if how large our efficiency gain is in practice, we have implemented this use case for ABY3 twice. One implementation is optimised as we described, the other implementation is fully unoptimized. . TODO naive implementation im appendix wenn zeit




Figure 6.2: Memory usage of ABY3, Conclave and SMCQL for our first use case

**Evaluation of Runtime** Our measurement shows that for small input sizes the unoptimised implementation is faster, while for larger input sizes the optimized implementation is faster. For a visualisation of our comparison see Figure 6.3. For input size 2000, the unoptimised implementation runs within 182 milliseconds, which is 1,5 times faster then our optimised implementation, that needs 276 milliseconds. With increasing input size the difference between their performances becomes less significant, for input size 8000 the unoptimized implementation is less then 1,4 times faster. For an input of size 64000 the two implementation have very similar runtime and both need about 500 milliseconds. From there on onwards, our optimization starts to pay of and is constantly about 1,5 times then the not optimised version. TODO ursache erklären warum erst langsamer dann schneller

**Evaluation of Space** From the perspective of memory consumption our optimization is a strict improvement over the unoptimized implementation. The optimized implementation is strict better, as for every single input size we have tested, it requires less memory. We visualise the memory requirements of our two implementations in Figure 6.4. It is notable how similar both implementations scale. As, for large inputs, they both very reliable double their memory consumption each time their input is doubled and the optimized implementation consistently takes half the memory of the unoptimised implementation. This trend start with input of size 64000 where our optimization needs about 300MB and the naive implementation about 150MB and contentious from thereon.




Figure 6.3: Runtime of our two implementations of ABY3 of use case two




Figure 6.4: Space requirement of our two implementations of ABY3 of use case two




Figure 6.5: Comparison between Conclave and ABY3 implementation of our second use case

**Comparison between Conclave and ABY3** In our second use case Conclave scales better than in the first one. For a visualisation of its results see figure 6.5. We have been able to obtain a result for 1000 input columns in 6 hours and 32 minutes. Compared to the 500 input columns in 8 hours from the first use case, Conclave has been able to handle an input twice as large in less time. Yet despite these better results Conclave is still unable, to compete with the results of our two ABY3 implementations for the second use case results, that both have been able to compute the result for 1000 input-rows in less then a second.

## 6.3 use case Four

As we described in Section 6.1 our fourth use case has an identical semantic to our first use case but we make use of Conclave’s trust annotations, that allow the leakage of some of the input data, which allows Conclave to apply optimizations that speed up the Computation. Therefore we focus on a comparison between the performance of Conclave in the first and forth use case. A visualisation of our comparison can found in Figure 6.6 and 6.7.

**Comparison between use case One and use case Four** With the usage of trust annotation we have observed an significant improvement in speed and efficiency. We have been able to evaluate the query for input-sizes of up to 3000 rows, which are 6 times larger then in use case one. For input-sizes larger then 3000 Conclave tends to




Figure 6.6: Runtime of Conclave in use case one and use case two

crash because of internal errors that are outside of our control. Therefore we are unable to obtain result for larger inputs. For an input of size 500 in use case four, Conclave needs less than 5 minutes which is more than 80 times faster than the 8 hours required in use case one. For an input of size 3000 Conclave is able to compute the correct result in less than 25 minutes. From the perspective of memory requirement the picture is very similar. As Conclave requires less memory in use case four than in use case one, for all input-sizes we observed.

**Comparison between ABY3 and Conclave** It is notable that even with the unfair advantage of leaking some input data, Conclave is not able to yield similar performance to ABY3. As our implementation of use case one with ABY3, that has no such advantage, still is multiple orders of magnitude faster than Conclave in use case four and also requires significant less memory. As on average it is more than 360 times faster. The difference in speed becomes even more significant with larger input-sizes, in the extreme case of input-size 3000 ABY3 is more than 1300 times faster than Conclave.




Figure 6.7: Comparison between Conclave's space consumption in use case one and use case four





# Bibliography

- [ABL<sup>+</sup>18] David W. Archer, Dan Bogdanov, Yehuda Lindell, Liina Kamm, Kurt Nielsen, Jakob Illeborg Pagter, Nigel P. Smart, and Rebecca N. Wright. From keys to databases - real-world applications of secure multi-party computation. *Comput. J.*, 61:1749–1771, 2018.
- [ACC<sup>+</sup>21] Abdelrahman Aly, K Cong, D Cozzo, M Keller, E Orsini, D Rotaru, O Scherer, P Scholl, N Smart, T Tanguy, et al. Scale-mamba v1. 12: Documentation, 2021.
- [AGJ<sup>+</sup>21] Surya Addanki, Kevin Garbe, Eli Jaffe, Rafail Ostrovsky, and Antigoni Polychroniadou. Prio+: Privacy preserving aggregate statistics via boolean shares. Cryptology ePrint Archive, Report 2021/576, 2021. <https://ia.cr/2021/576>.
- [AIL<sup>+</sup>19] Kinan Dak Albab, Rawane Issa, Andrei Lapets, Peter Flockhart, Lucy Qin, and Ira Globus-Harris. Tutorial: Deploying secure multi-party computation on the web using jiff. *2019 IEEE Cybersecurity Development (SecDev)*, pages 3–3, 2019.
- [BEE<sup>+</sup>16] Johes Bater, Gregory Elliott, Craig Eggen, Satyender Goel, Abel Kho, and Jennie Rogers. Smcql: Secure querying for federated databases. *arXiv preprint arXiv:1606.06808*, 2016.
- [BJSV15] Dan Bogdanov, Marko Jõemets, Sander Siim, and Meril Vaht. How the estonian tax and customs board evaluated a tax fraud detection system based on secure multi-party computation. In *International conference on financial cryptography and data security*, pages 227–234. Springer, 2015.
- [CD05] Ronald Cramer and Ivan Damgård. Multiparty computation, an introduction. In *Contemporary cryptology*, pages 41–87. Springer, 2005.
- [CGB17] Henry Corrigan-Gibbs and Dan Boneh. Prio: Private, robust, and scalable computation of aggregate statistics. In *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*, pages 259–282, Boston, MA, March 2017. USENIX Association.
- [Cos21] Cosmian. Ciphercompute. <https://github.com/Cosmian/CipherCompute>, 2021.

- [cPr21] The python profilers. <https://docs.python.org/3/library/profile.html>, 2021.
- [DPSZ11] I. Damgard, V. Pastro, N. P. Smart, and S. Zakarias. Multiparty computation from somewhat homomorphic encryption. Cryptology ePrint Archive, Paper 2011/535, 2011. <https://eprint.iacr.org/2011/535>.
- [DW82] Danny Dolev and Avi Wigderson. On the security of multi-party protocols in distributed systems. In David Chaum, Ronald L. Rivest, and Alan T. Sherman, editors, *Advances in Cryptology: Proceedings of CRYPTO '82, Santa Barbara, California, USA, August 23-25, 1982*, pages 167–175. Plenum Press, New York, 1982.
- [HHNZ19] Marcella Hastings, Brett Hemenway, Daniel Noble, and Steve Zdancewic. Sok: General purpose compilers for secure multi-party computation. In *2019 IEEE symposium on security and privacy (SP)*, pages 1220–1237. IEEE, 2019.
- [JNS17] Noah M Johnson, Joseph P Near, and Dawn Xiaodong Song. Practical differential privacy for sql queries using elastic sensitivity. *CoRR*, abs/1706.09479, 2017.
- [Lin17] Yehuda Lindell. How to simulate it—a tutorial on the simulation proof technique. *Tutorials on the Foundations of Cryptography*, pages 277–346, 2017.
- [LWN<sup>+</sup>15] Chang Liu, Xiao Shaun Wang, Kartik Nayak, Yan Huang, and Elaine Shi. Oblivm: A programming framework for secure computation. In *2015 IEEE Symposium on Security and Privacy*, pages 359–376. IEEE, 2015.
- [PSZ14] Benny Pinkas, Thomas Schneider, and Michael Zohner. Faster private set intersection based on  $\{\text{OT}\}$  extension. In *23rd USENIX Security Symposium (USENIX Security 14)*, pages 797–812, 2014.
- [RAB<sup>+</sup>22] Jennie Rogers, Elizabeth Adetoro, Johes Bater, Talia Canter, Dong Fu, Andrew Hamilton, Amro Hassan, Ashley Martinez, Erick Michalski, Vesna Mitrovic, et al. Vaultdb: A real-world pilot of secure multi-party computation within a clinical research network. *arXiv preprint arXiv:2203.00146*, 2022.
- [Rin] Peter Rindal. The ABY3 Framework for Machine Learning and Database Operations. <https://github.com/ladnir/aby3>.
- [Rin21] Peter Rindal. cryptotools. <https://github.com/ladnir/cryptoTools>, 2021.
- [Sho22] Shopify. toxiproxy. <https://github.com/Shopify/toxiproxy>, 2022.

- [tim21a] time. <https://manpages.debian.org/stretch/time/time.1.en.html>, 2021.
- [tim21b] timeit. <https://docs.python.org/3/library/timeit.html>, 2021.
- [VSG<sup>+</sup>19] Nikolaj Volgushev, Malte Schwarzkopf, Ben Getchell, Mayank Varia, Andrei Lapets, and Azer Bestavros. Conclave: Secure multi-party computation on big data. In *Proceedings of the Fourteenth EuroSys Conference 2019*, EuroSys '19, New York, NY, USA, 2019. Association for Computing Machinery.
- [wir22] wireshark. <https://github.com/multiparty/jiff>, 2022.
- [WMK16] Xiao Wang, Alex J. Malozemoff, and Jonathan Katz. EMP-toolkit: Efficient MultiParty computation toolkit. <https://github.com/emp-toolkit>, 2016.
- [Yao82] Andrew C. Yao. Protocols for secure computations. In *23rd Annual Symposium on Foundations of Computer Science (sfcs 1982)*, pages 160–164, 1982.